

Machine Learning techniques for numerical solvers

Day 2: supervised learning

Maria Han Veiga
Mini-course SUSTech
11.03 - 14.03



THE OHIO STATE UNIVERSITY

COLLEGE OF ARTS AND SCIENCES

Schedule

- **Monday:** Introduction to Machine Learning (1h) 17:00-18:00
- **Tuesday:** Computational Framework + Supervised learning: integrating data-driven methods within a numerical solver + Hands-on session (2h) 14:00-16:00
- **Wednesday:** Unsupervised learning: Physics informed neural networks (1h) 11:00-12:00
- **Thursday:** Reinforcement Learning ? (1h) 11:00-12:00

Day 3: Unsupervised learning

Outline

- Physics informed neural networks (PINNs)
 - Raissi et al. 2017
 - PINNs now
- Hands-on session:
 - Viscous Burgers equation in 1D
 - ODE — damped oscillator

Motivation

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .

Motivation

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .

- **Idea:** Use a data-driven method to solve the PDE problem.
 - Dissanayake et al, 1994, Lagalis et al. 1998 — Artificial neural networks to solve ODEs/PDEs
 - Raissi et al. 2017, Sirignano et al. 2017 — Novel minimisation problem and modern deep learning techniques

Physics informed neural networks (PINNs)

- According to authors:
- Neural networks trained to “*solve supervised learning tasks while respecting physical laws (PDEs)*” (Raissi et al. 2017)
 - Data-driven solution
- Two types of algorithms: (Raissi et al. 2019)
 - Family of data-efficient spatio-temporal function approximators
 - Arbitrary accurate RK time steppers

Physics informed neural networks (PINNs)

- According to authors:
- Neural networks trained to “*solve supervised learning tasks while respecting physical laws (PDEs)*” (Raissi et al. 2017)
 - Data-driven solution
- Two types of algorithms: (Raissi et al. 2019)
 - Family of data-efficient spatio-temporal function approximators
 - Arbitrary accurate RK time steppers
- I call these unsupervised learning, or maybe self-supervised learning, because we don't have explicitly a dataset $S = \{(x_i, y_i), i = 1, \dots, m\}$.

Setup

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .

Setup

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .

e.g. Viscous Burger's equation in 2D:

$$\mathcal{L}(u; \lambda) = \lambda_1 u \partial_x u - \lambda_2 u \partial_{xx} u$$

Setup

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .

Task:

- 1) Given λ , what is $u(x, t)$ that fulfils the PDE? (Data-driven solution of PDEs)
- 2) Find λ that best describes observations $u(x_i, t_i)$. (Data-driven discovery of PDEs)

Setup

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .

Task:

- 1) **Given λ , what is $u(x, t)$ that fulfils the PDE? (Data-driven solution of PDEs)**
- 2) Find λ that best describes observations $u(x_i, t_i)$. (Data-driven discovery of PDEs)

Data-driven solutions

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

$$u(x, t = 0) = u_0(x)$$

$$u(\partial\Omega, t) = g(t)$$

Data-driven solutions

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

$$u(x, t = 0) = u_0(x)$$

$$u(\partial\Omega, t) = g(t)$$

- Let the solution of the PDE be given by a neural network $u(t, x) = \mathcal{NN}(t, x; \theta)$

Data-driven solutions

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

$$u(x, t = 0) = u_0(x)$$

$$u(\partial\Omega, t) = g(t)$$

- Let the solution of the PDE be given by a neural network $u(t, x) = \mathcal{NN}(t, x; \theta)$
- To solve the PDE, we want the residual $f(u; t, x) := \partial_t u + \mathcal{L}(u)$ to be zero at all (x, t) and that the neural network fulfils the initial and boundary data.

Data-driven solutions

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

$$u(x, t = 0) = u_0(x)$$

$$u(\partial\Omega, t) = g(t)$$

- Let the solution of the PDE be given by a neural network $u(t, x) = \mathcal{NN}(t, x; \theta)$
- To solve the PDE, we want the residual $f(u; t, x) := \partial_t u + \mathcal{L}(u)$ to be zero at all (x, t) and that the neural network fulfils the initial and boundary data.
- Introduce loss:

$$L(\mathcal{NN}) = \frac{1}{N_i} \lambda_i \sum_{i=1}^{N_i} |\mathcal{NN}(x_i, 0) - u(x_i, 0)|^2 + \frac{1}{N_b} \lambda_b \sum_{i=1}^{N_b} |\mathcal{NN}(\partial\Omega_i, t_i) - g(x_i)|^2 \\ + \frac{1}{N_f} \lambda_f \sum_{i=1}^{N_f} |f(x_i, t_i)|^2$$

Data-driven solutions

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

$$u(x, t = 0) = u_0(x)$$

$$u(\partial\Omega, t) = g(t)$$

- Let the solution of the PDE be given by a neural network $u(t, x) = \mathcal{NN}(t, x; \theta)$
- To solve the PDE, we want the residual $f(u; t, x) := \partial_t u + \mathcal{L}(u)$ to be zero at all (x, t) and that the neural network fulfils the initial and boundary data.
- Introduce loss:

$$L(\mathcal{NN}) = \frac{1}{N_i} \lambda_i \sum_{i=1}^{N_i} |\mathcal{NN}(x_i, 0) - u(x_i, 0)|^2 + \frac{1}{N_b} \lambda_b \sum_{i=1}^{N_b} |\mathcal{NN}(\partial\Omega_i, t_i) - g(x_i)|^2 \\ + \frac{1}{N_f} \lambda_f \sum_{i=1}^{N_f} |f(x_i, t_i)|^2$$

Data-driven solutions

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

$$u(x, t = 0) = u_0(x)$$

$$u(\partial\Omega, t) = g(t)$$

- Let the solution of the PDE be given by a neural network $u(t, x) = \mathcal{NN}(t, x; \theta)$
- To solve the PDE, we want the residual $f(u; t, x) := \partial_t u + \mathcal{L}(u)$ to be zero at all (x, t) and that the neural network fulfils the initial and boundary data.
- Introduce loss:

$$L(\mathcal{NN}) = \frac{1}{N_i} \lambda_i \sum_{i=1}^{N_i} |\mathcal{NN}(x_i, 0) - u(x_i, 0)|^2 + \frac{1}{N_b} \lambda_b \sum_{i=1}^{N_b} |\mathcal{NN}(\partial\Omega_i, t_i) - g(x_i)|^2$$

$$+ \frac{1}{N_f} \lambda_f \sum_{i=1}^{N_f} |f(x_i, t_i)|^2$$

Data-driven solutions

$$L(\mathcal{NN}) = \frac{1}{N_i} \lambda_i \sum_{i=1}^{N_i} | \mathcal{NN}(x_i, 0) - u(x_i, 0) |^2 + \frac{1}{N_b} \lambda_b \sum_{i=1}^{N_b} | \mathcal{NN}(\partial\Omega_i, t_i) - g(x_i) |^2 \\ + \frac{1}{N_f} \lambda_f \sum_{i=1}^{N_f} | f(x_i, t_i) |^2$$

Data-driven solutions

$$L(\mathcal{NN}) = \frac{1}{N_i} \lambda_i \sum_{i=1}^{N_i} | \mathcal{NN}(x_i, 0) - u(x_i, 0) |^2 + \frac{1}{N_b} \lambda_b \sum_{i=1}^{N_b} | \mathcal{NN}(\partial\Omega_i, t_i) - g(x_i) |^2 \\ + \frac{1}{N_f} \lambda_f \sum_{i=1}^{N_f} | f(x_i, t_i) |^2$$

The loss is composed by two types of nodes:

- Initial and boundary nodes
- Collocation points

Data-driven solutions

- No theoretical guarantees but if PDE is well-posed, seems like optimiser will find a solution.

Data-driven solutions

- No theoretical guarantees but if PDE is well-posed, seems like optimiser will find a solution.
- Example from PINNs paper (Raissi et al 2017), solve the 1D Burgers' equation with Dirichlet boundary conditions

$$\partial_t u + u \partial_x u - (0.01/\pi) \partial_{xx} u = 0 \quad x \in [-1, 1], \quad t \in [0, 1]$$

$$u(x, 0) = -\sin(\pi x)$$

$$u(-1, t) = u(1, t) = 0$$

Data-driven solutions

- No theoretical guarantees but if PDE is well-posed, seems like optimiser will find a solution.
- Example from PINNs paper (Raissi et al 2017), solve the 1D Burgers' equation with Dirichlet boundary conditions

$$\partial_t u + u \partial_x u - (0.01/\pi) \partial_{xx} u = 0 \quad x \in [-1, 1], \quad t \in [0, 1]$$

$$u(x, 0) = -\sin(\pi x)$$

$$u(-1, t) = u(1, t) = 0$$


$$\text{Define } f := \partial_t u + u \partial_x u - (0.01/\pi) \partial_{xx} u.$$

Let $u(x, t) = \mathcal{N}\mathcal{N}(x, t)$ and minimise Physics Informed Loss:

$$L(\mathcal{N}\mathcal{N}) = \text{Initial data loss} + \text{Boundary data loss} + \text{Residual loss}$$

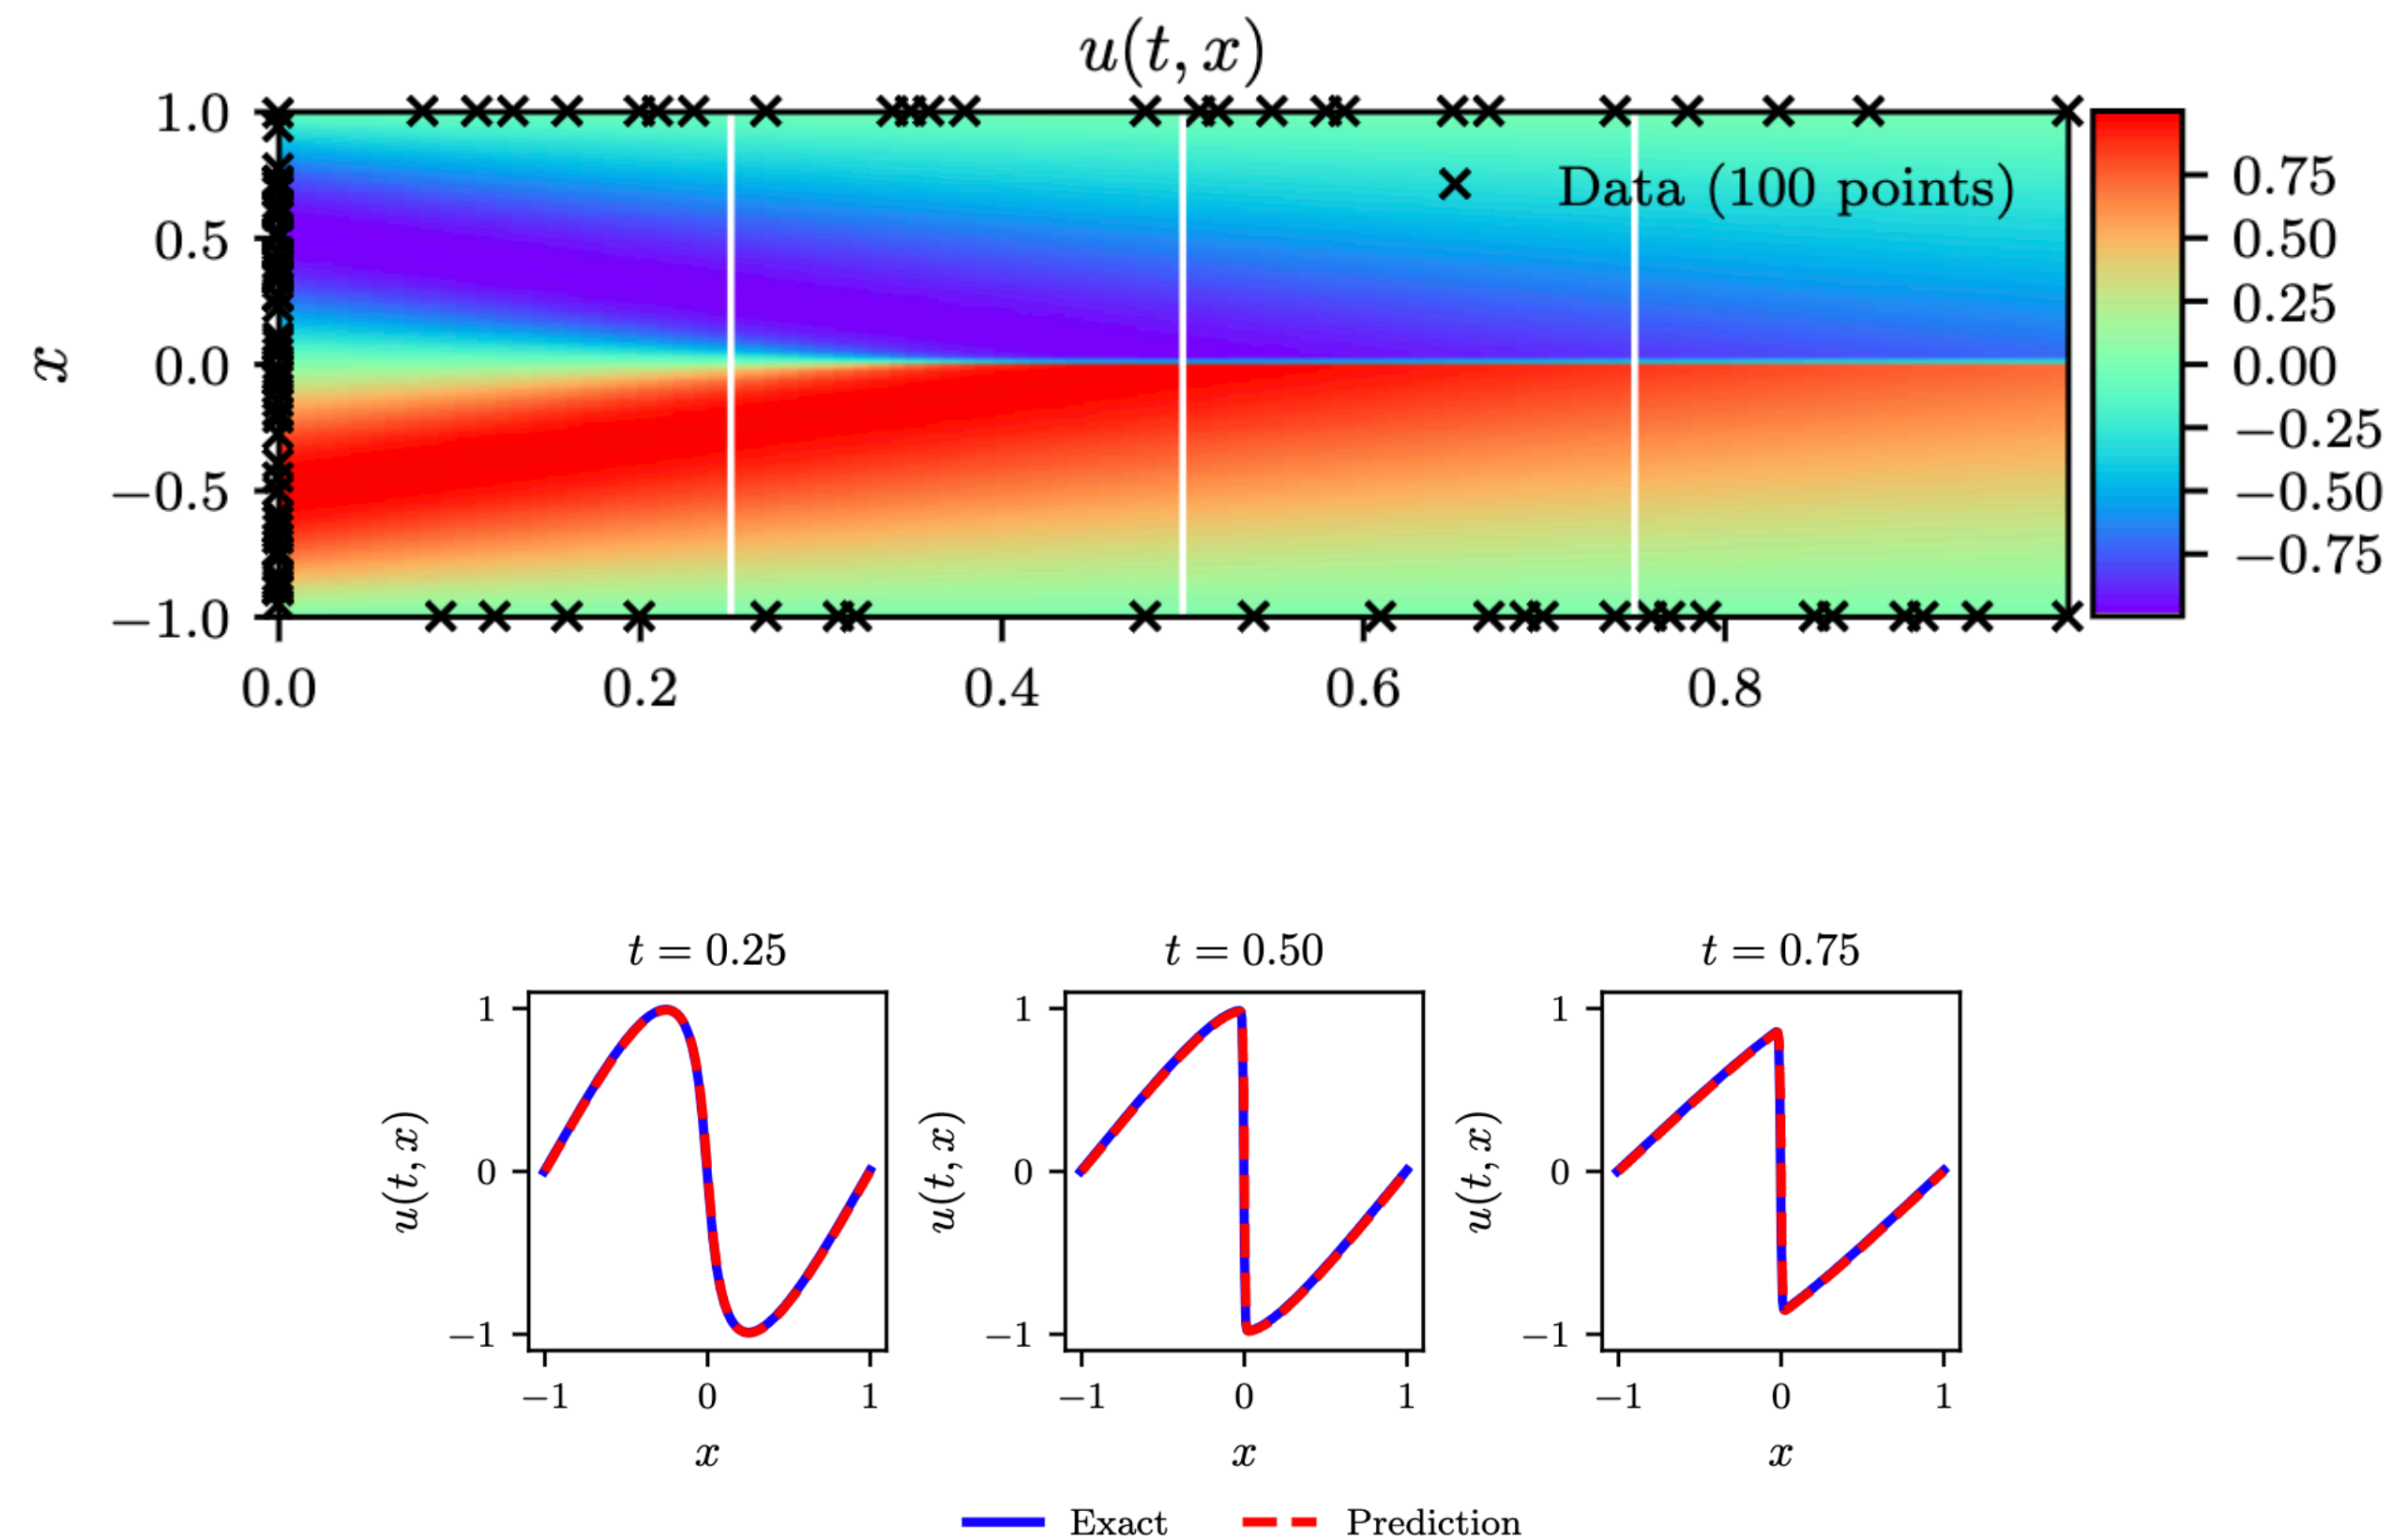
Data-driven solutions

- Example from PINNs paper (Raissi et al 2017)



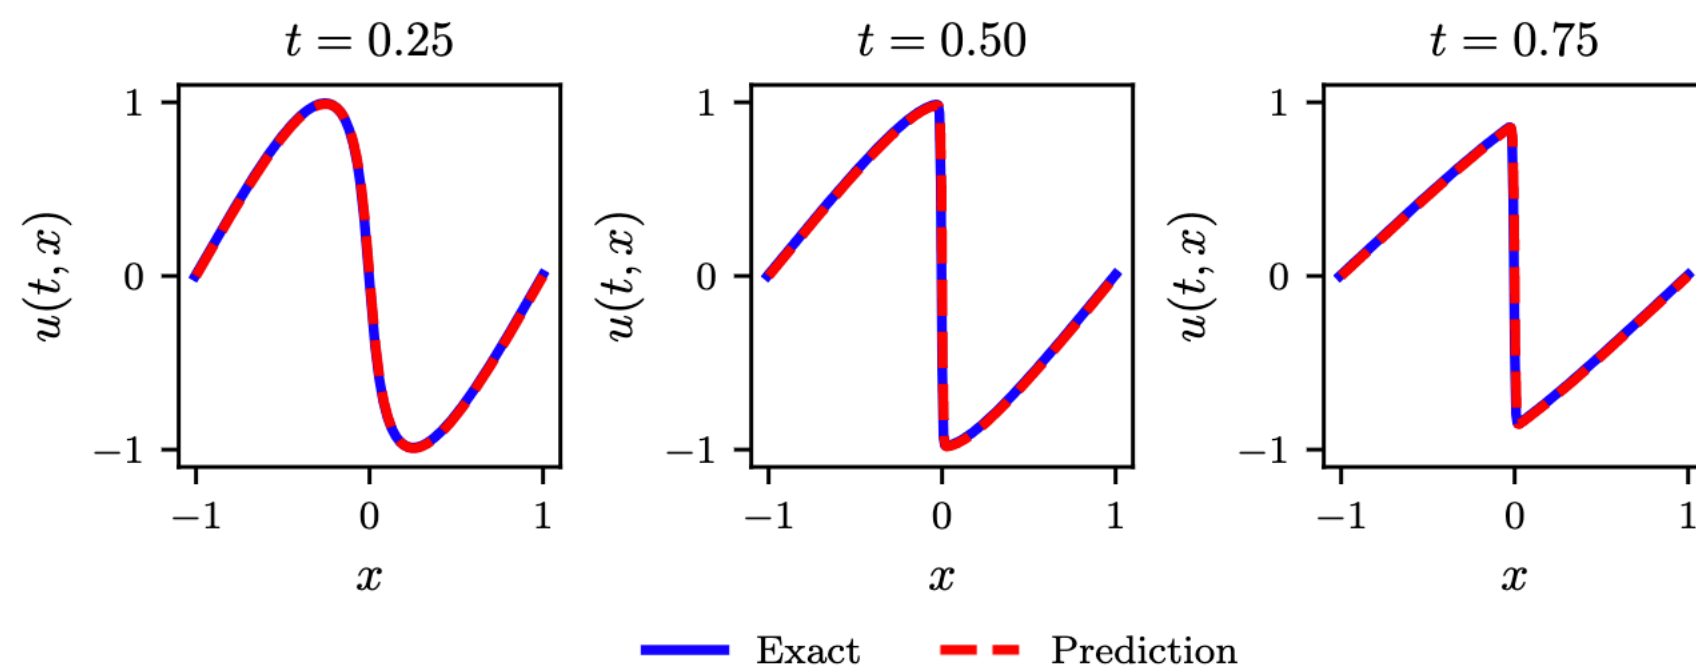

Data-driven solutions

- Example from PINNs paper (Raissi et al 2017)



Data-driven solutions

- Example from PINNs paper (Raissi et al 2017)



- Training details:
 - 100 points for boundary and initial data
 - 10000 colocation points (for residual), sampled using Latin Hypercube Sampling strategy. (2D space)

Data-driven solutions

- In original PINNs paper (Raissi et al 2017)
 - Shrödinger equation in 1D (complex numbers)
 - Allen-Cahn Equation

Pros and cons

Pros

- Mesh free
- Mostly unsupervised
- Can work with noisy data

Cons:

- Convergence properties are not well understood
- Computational (training) cost can be much higher than a traditional solver
- Poor scaling to large domains / high frequencies / more complex solutions

PINNs now

- Review paper (Cuomo et al 2022):
 - Different types of PINNs: beyond fully connected
 - CNNs, AutoEncoders, ResNet, Recurrent (often tailor made for different applications)
 - Different type of PDEs (see references in review paper):
 - Advection-diffusion-reaction problems
 - Diffusion problems
 - Advection problems
 - Navier-Stokes equations
 - Hyperbolic equations

<https://link.springer.com/article/10.1007/s10915-022-01939-z>

PINNs now

- Review paper (Cuomo et al 2022):
 - Different types of PINNs: beyond fully connected
 - CNNs, AutoEncoders, ResNet, Recurrent (often tailor made for different applications)
 - Different type of PDEs (see references in review paper):
 - Advection-diffusion-reaction problems
 - Diffusion problems
 - Advection problems
 - Navier-Stokes equations
 - **Hyperbolic equations**


<https://link.springer.com/article/10.1007/s10915-022-01939-z>

PINNs now

- PINNs for Hyperbolic equations:
 - Mao et al. 2019: 1D Euler equation

PINNs now


- PINNs for Hyperbolic equations:
 - Mao et al. 2019: 1D Euler equation
 - Example 1: Moving contact discontinuity, with shock position at $x = 0.5$



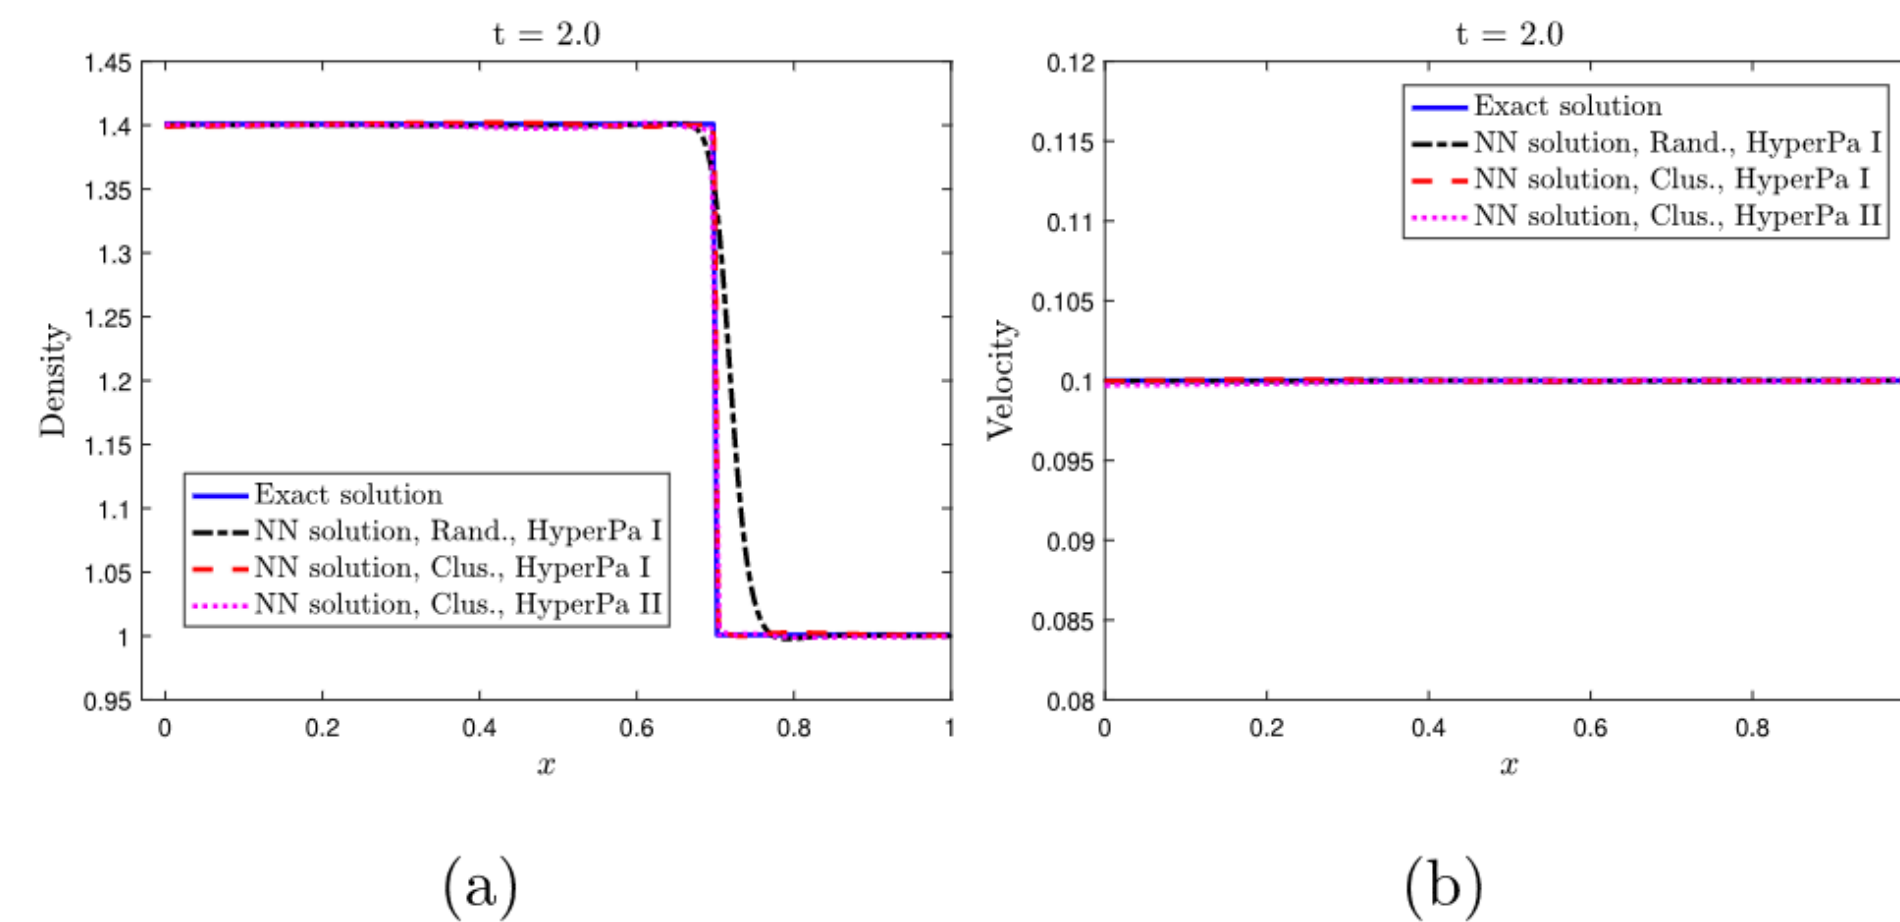
- Training nodes: 1120 points

PINNs now

- PINNs for Hyperbolic equations:
 - Mao et al. 2019: 1D Euler equation
 - Example 1: Moving contact discontinuity, with shock position at $x = 0.5$



- Training nodes: 1120 points

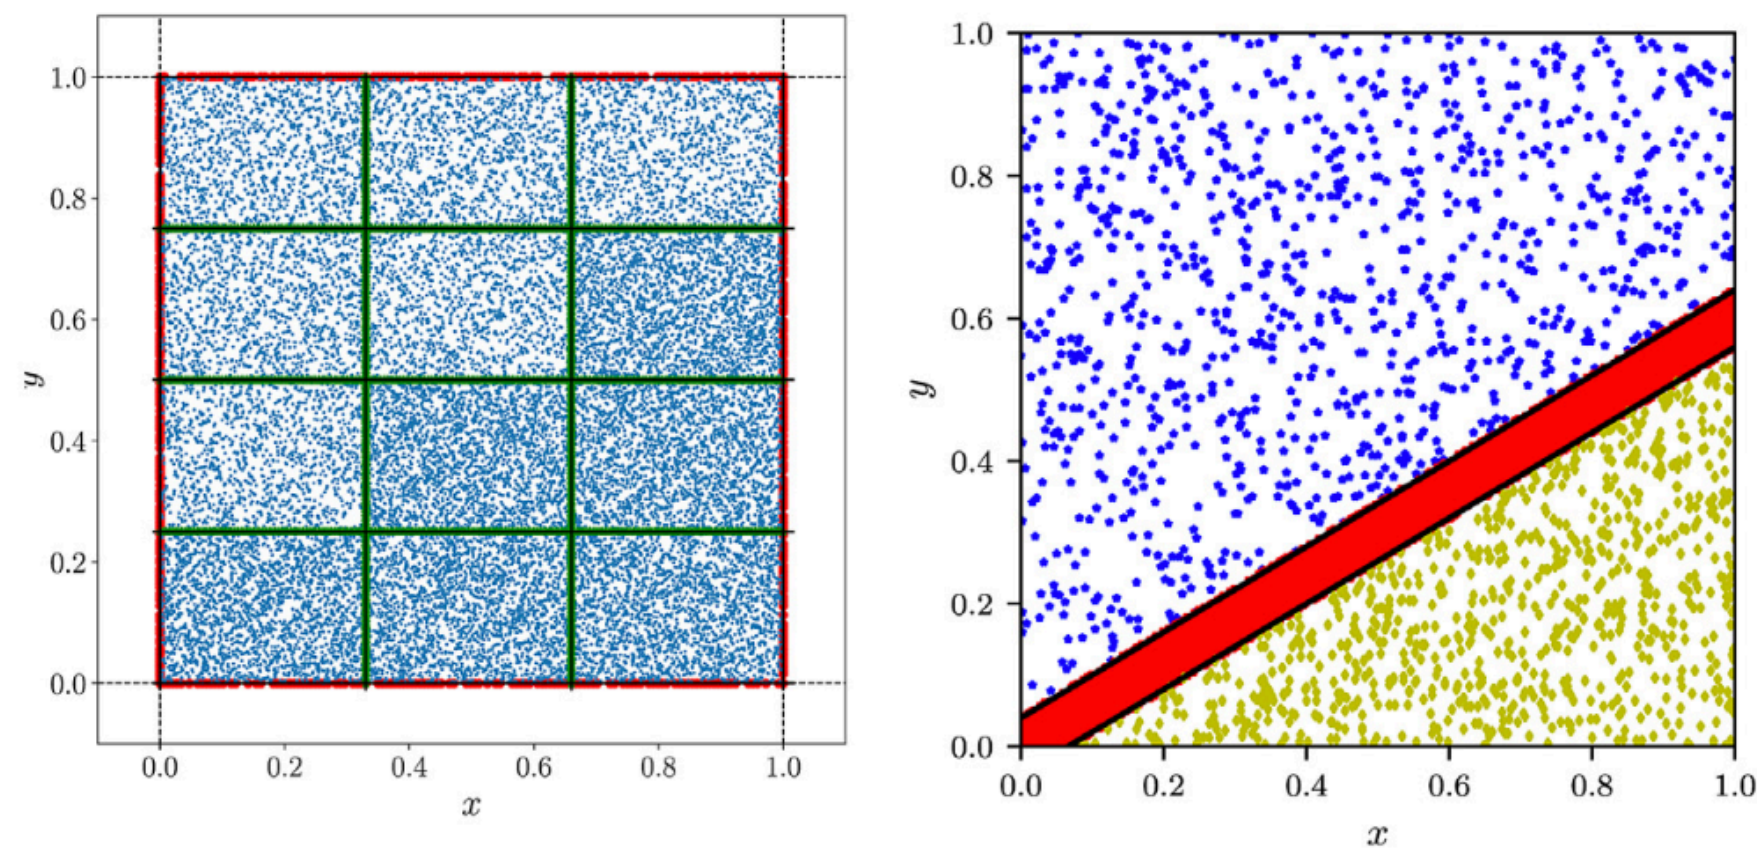


PINNs now

- PINNs for Hyperbolic equations:
 - Jagtap et al. 2020: 1D and 2D Euler equation
 - PINNs loss includes also conservation terms and domain decomposition
 - Example 2D Euler, IC: oblique shock wave problem

PINNs now

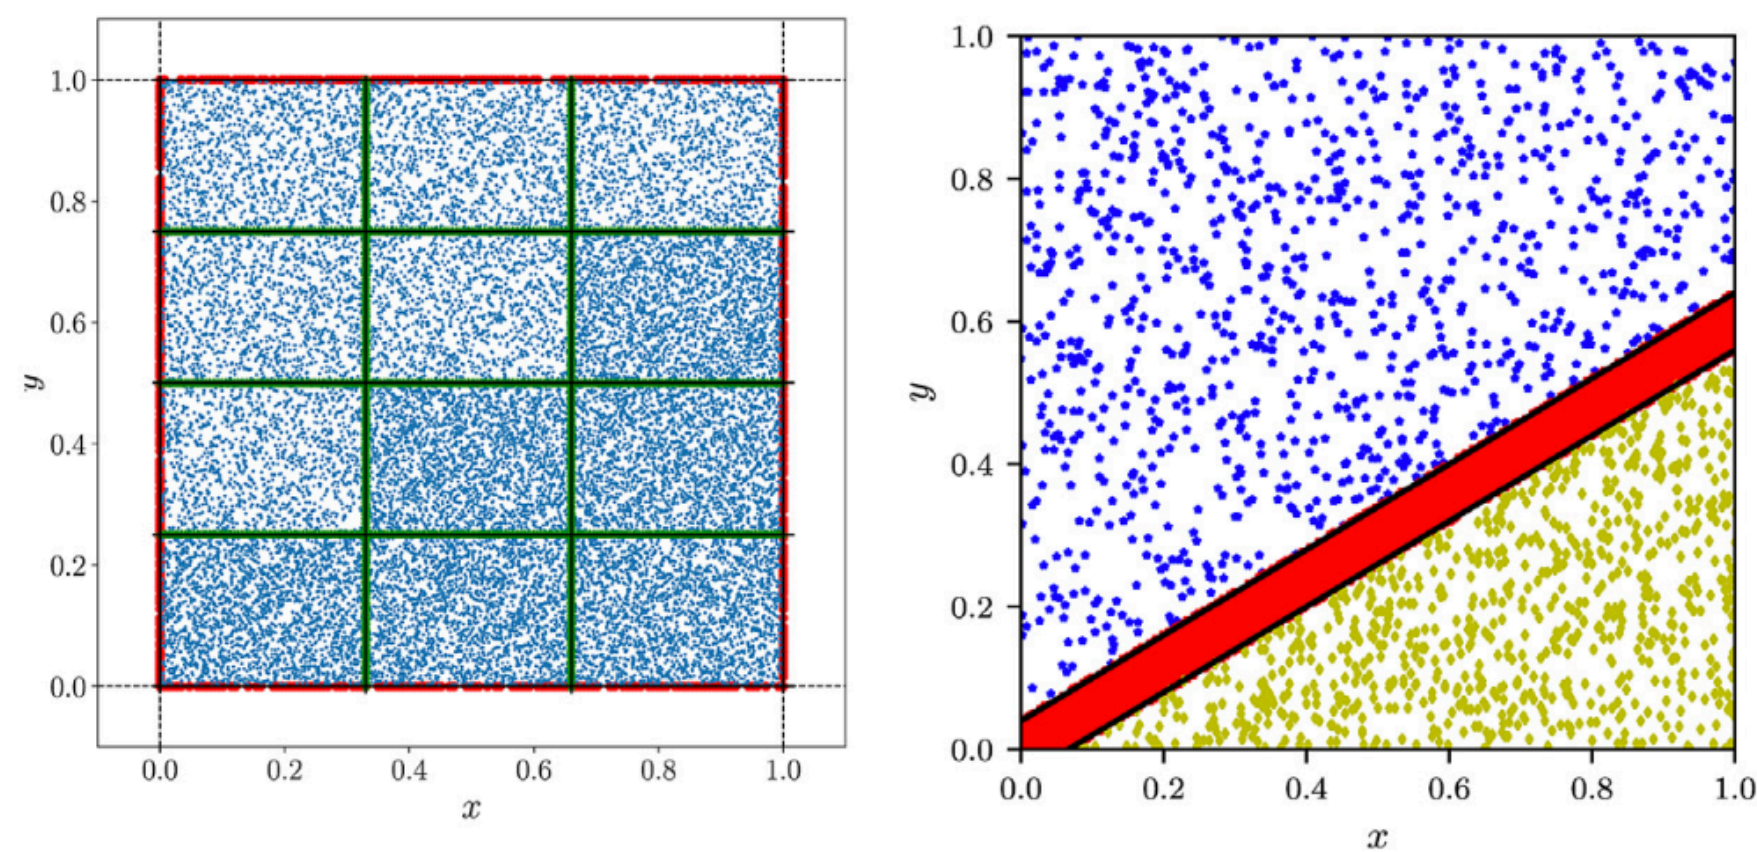
- PINNs for Hyperbolic equations:
 - Jagtap et al. 2020: 1D and 2D Euler equation
 - PINNs loss includes also conservation terms and domain decomposition
 - Example 2D Euler, IC: oblique shock wave problem



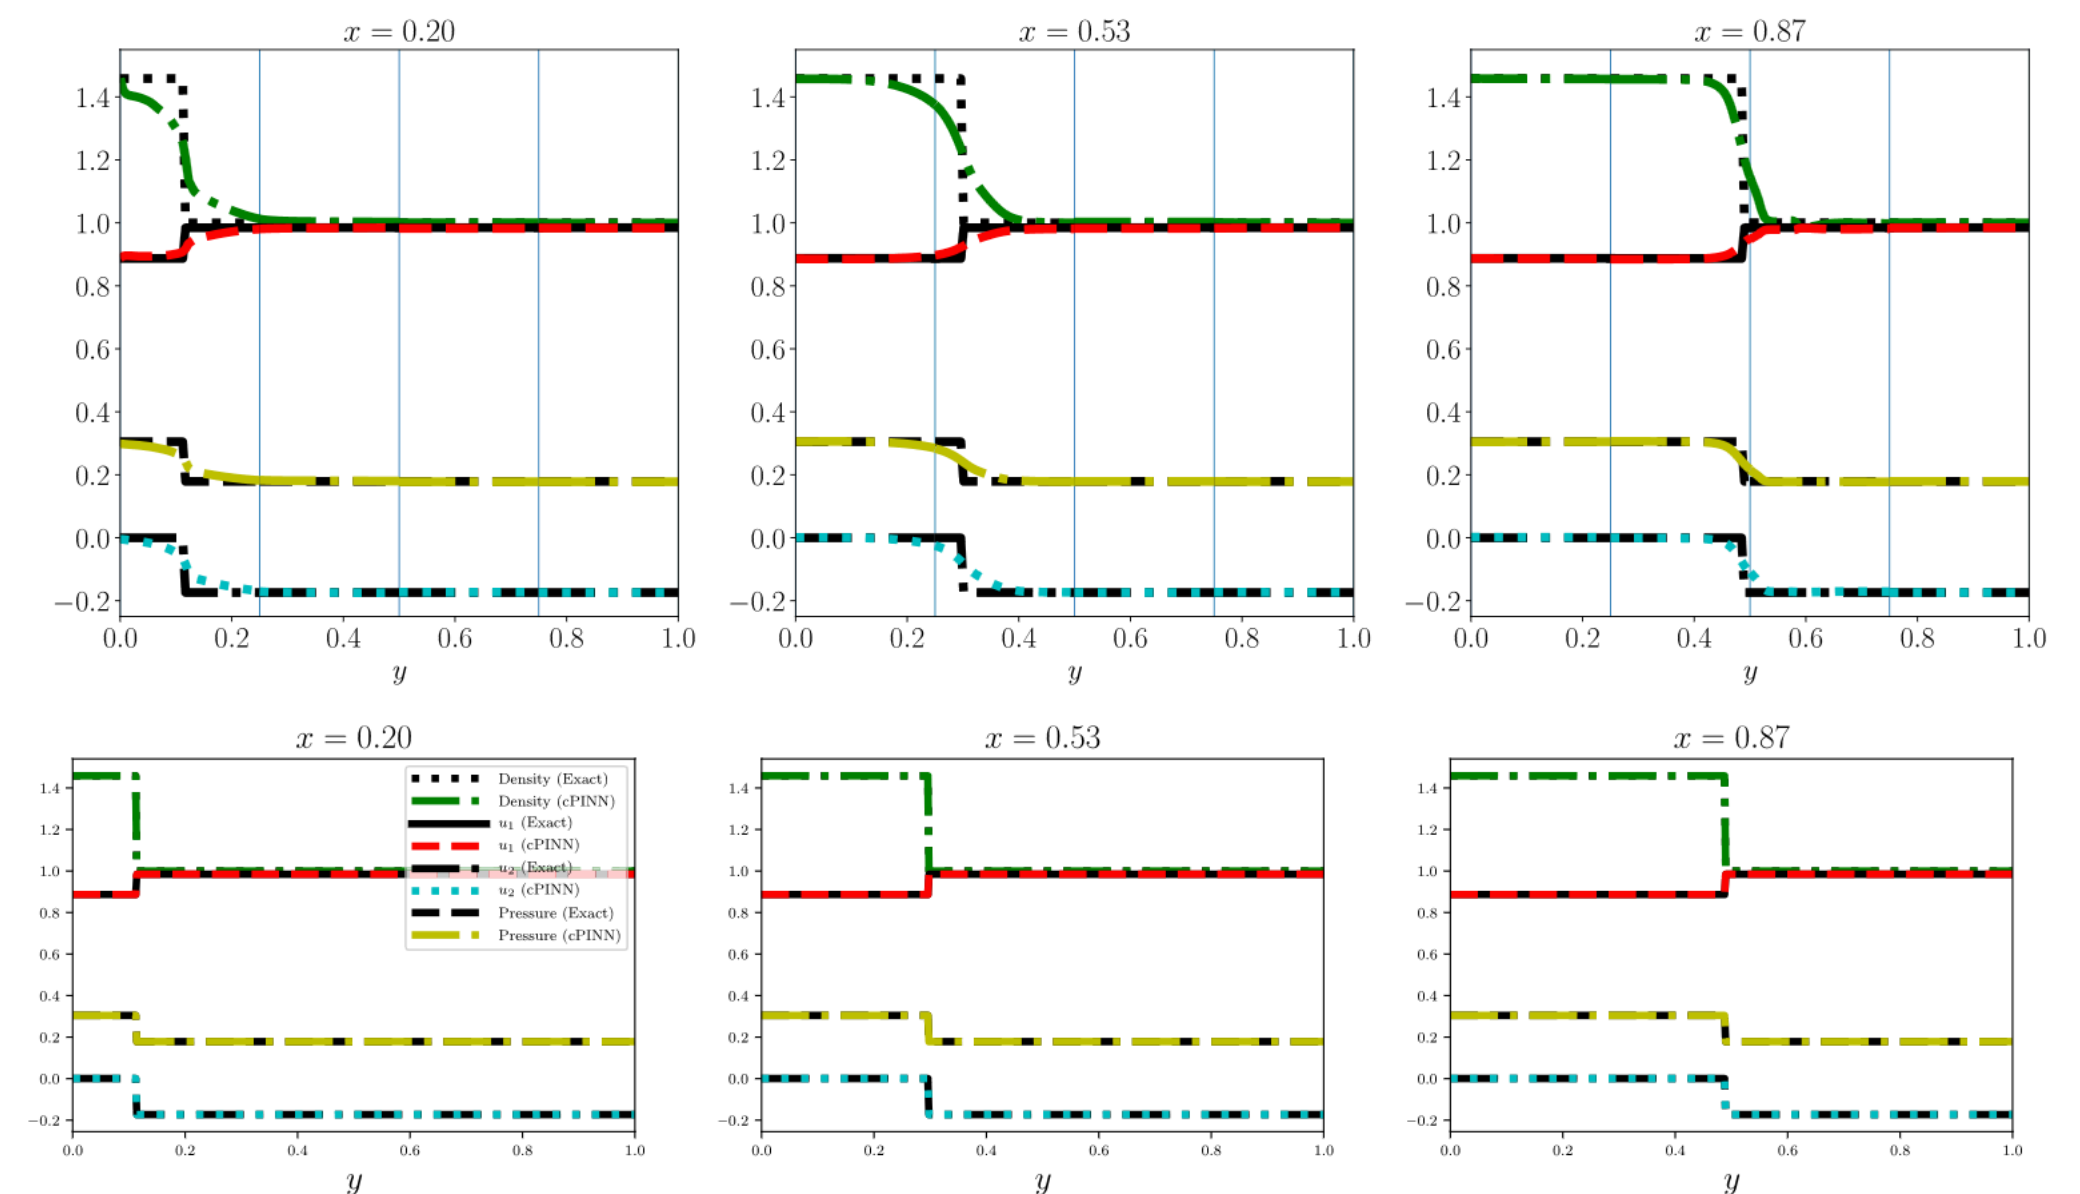
Domain decomposition

PINNs now

- PINNs for Hyperbolic equations:
 - Jagtap et al. 2020: 1D and 2D Euler equation
 - PINNs loss includes also conservation terms and domain decomposition
 - Example 2D Euler, IC: oblique shock wave problem



Domain decomposition



PINNs now

- Error analysis (Kutyniok, 2022).
- The global error between a trained deep neural network u_{θ}^* and the correct solution function u for a PDE can be bounded by:

$$\mathcal{R}(u_{\theta}^*) \leq \mathcal{E}_0 + 2\mathcal{E}_G + \mathcal{E}_A$$

PINNs now

- Error analysis (Kutyniok, 2022).
- The global error between a trained deep neural network u_{θ}^* and the correct solution function u for a PDE can be bounded by:

$$\mathcal{R}(u_{\theta}^*) \leq \mathcal{E}_0 + 2\mathcal{E}_G + \mathcal{E}_A$$

- \mathcal{E}_0 is the *optimisation error*, accounting for GD not finding the absolute minimum of the loss

PINNs now

- Error analysis (Kutyniok, 2022).
- The global error between a trained deep neural network u_{θ}^* and the correct solution function u for a PDE can be bounded by:

$$\mathcal{R}(u_{\theta}^*) \leq \mathcal{E}_0 + 2\mathcal{E}_G + \mathcal{E}_A$$

- \mathcal{E}_0 is the *optimisation error*, accounting for GD not finding the absolute minimum of the loss
- \mathcal{E}_g is the *generalisation gap* (Mishra et al. 2022)

PINNs now

- Error analysis (Kutyniok, 2022).
- The global error between a trained deep neural network u_{θ}^* and the correct solution function u for a PDE can be bounded by:

$$\mathcal{R}(u_{\theta}^*) \leq \mathcal{E}_0 + 2\mathcal{E}_G + \mathcal{E}_A$$

- \mathcal{E}_0 is the *optimisation error*, accounting for GD not finding the absolute minimum of the loss
- \mathcal{E}_g is the *generalisation gap* (Mishra et al. 2022)
- \mathcal{E}_A is the *approximation error*, the ability of the neural network to approximate the exact solution (De Ryck et al. 2021)

PINNs inverse problems

- So far we focused on solving the *forward problem*, however, it seems like PINNs is quite effective to solve inverse problems / data-driven discovery of PDEs.

PINNs inverse problems

- So far we focused on solving the *forward problem*, however, it seems like PINNs is quite effective to solve inverse problems / data-driven discovery of PDEs.

- Parametrised, nonlinear PDE(s)

$$\partial_t u + \mathcal{L}(u; \lambda) = 0, \quad x \in \Omega \subset \mathbb{R}^D, \quad t \in [0, T]$$

- where $u(x, t)$ denotes the solution and $\mathcal{L}(\cdot; \lambda)$ is a nonlinear operator parametrised by λ .
- One has observations of the solution $u(x, t)$, find the correct λ
- $\mathcal{L}(u, \cdot) := (\cdot) \partial_x u + (\cdot) \partial_{xx} u + (\cdot) \partial_y u + (\cdot) \partial_{yy} u + \dots$

Hands-on

PINNs for 1D viscous burgers and ODE

PDE

- Solve the 1D Burgers' equation with Dirichlet boundary conditions

$$\partial_t u + u \partial_x u - (0.01/\pi) \partial_{xx} u = 0 \quad x \in [-1, 1], \quad t \in [0, 1]$$

$$u(x, 0) = -\sin(\pi x)$$

$$u(-1, t) = u(1, t) = 0$$

PDE

Define the sampling nodes

```
: boundary_nodes = np.zeros([100,2])  
boundary_nodes[0:50,0]=1  
boundary_nodes[0:50,1]=np.random.uniform(low=0.0, high=1, size=(50,))  
boundary_nodes[50:,0]=-1  
boundary_nodes[50:,1]=np.random.uniform(low=0.0, high=1, size=(50,))
```

```
: initial_data_nodes = np.zeros([100,2])  
initial_data_nodes[:,0]=np.random.uniform(low=-1.0, high=1, size=(100,))  
initial_data_nodes[:,1]=0  
initial_data_values = -np.sin(np.pi*initial_data_nodes[:,0])
```

```
: collocation_data_nodes = np.zeros([5000,2])  
collocation_data_nodes[:,0]=np.random.uniform(low=-1.0, high=1, size=(5000,))  
collocation_data_nodes[:,1]=np.random.uniform(low=0.0, high=1, size=(5000,))
```

PDE

Define the sampling nodes

```
: boundary_nodes = np.zeros([100,2])  
boundary_nodes[0:50,0]=1  
boundary_nodes[0:50,1]=np.random.uniform(low=0.0, high=1, size=(50,))  
boundary_nodes[50:,0]=-1  
boundary_nodes[50:,1]=np.random.uniform(low=0.0, high=1, size=(50,))
```

```
: initial_data_nodes = np.zeros([100,2])  
initial_data_nodes[:,0]=np.random.uniform(low=-1.0, high=1, size=(100,))  
initial_data_nodes[:,1]=0  
initial_data_values = -np.sin(np.pi*initial_data_nodes[:,0]) ←  $u(x, t = 0) = -\sin(\pi x)$ 
```

```
: collocation_data_nodes = np.zeros([5000,2])  
collocation_data_nodes[:,0]=np.random.uniform(low=-1.0, high=1, size=(5000,))  
collocation_data_nodes[:,1]=np.random.uniform(low=0.0, high=1, size=(5000,))
```

PDE

Build the neural network:

- 2-dimensional input (x, t) , 1-dimensional output u
- 3 hidden layers
- 32 width in each hidden layer
- Activation: Tanh — important because it has to be differentiable!

PDE

Training cycle:

```
for i in range(5000):
    optimiser.zero_grad()

    lambda1, lambda2, lambda3 = 1, 1e-4, 1e-1

    # compute initial data loss
    u = pinn(initial_data_nodes_t)
    loss1 = torch.mean((torch.squeeze(u) - initial_data_values_t)**2)

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss2 = torch.mean((u - torch.zeros(u.shape))**2)

    # compute physics loss
    u = pinn(collocation_data_nodes_t)
    du = torch.autograd.grad(u, collocation_data_nodes_t, \
                             torch.ones([collocation_data_nodes_t.shape[0], 1]), \
                             retain_graph=True, create_graph=True)[0]
    du2 = torch.autograd.grad(du, collocation_data_nodes_t, \
                              torch.ones(collocation_data_nodes_t.shape), \
                              create_graph=True)[0]

    u_x = du[:, [0]]
    u_t = du[:, [1]]
    u_xx = du2[:, [0]]
    loss3 = torch.mean((u_t + u*u_x - .01 / np.pi *u_xx)**2)

    # backpropagate joint loss, take optimiser step
    loss = lambda1*loss1 + lambda2*loss2 + lambda3*loss3
    loss.backward()
    optimiser.step()
```


PDE

Training cycle:

$$|\mathcal{N}\mathcal{N}(x_i) - u_0(x_i)| \longrightarrow$$

```
for i in range(5000):
    optimiser.zero_grad()

    lambda1, lambda2, lambda3 = 1, 1e-4, 1e-1

    # compute initial data loss
    u = pinn(initial_data_nodes_t)
    loss1 = torch.mean((torch.squeeze(u) - initial_data_values_t)**2)

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss2 = torch.mean((u - torch.zeros(u.shape))**2)

    # compute physics loss
    u = pinn(collocation_data_nodes_t)
    du = torch.autograd.grad(u, collocation_data_nodes_t, \
                              torch.ones([collocation_data_nodes_t.shape[0], 1]), \
                              retain_graph=True, create_graph=True)[0]
    du2 = torch.autograd.grad(du, collocation_data_nodes_t, \
                              torch.ones(collocation_data_nodes_t.shape), \
                              create_graph=True)[0]

    u_x = du[:, [0]]
    u_t = du[:, [1]]
    u_xx = du2[:, [0]]
    loss3 = torch.mean((u_t + u*u_x - .01 / np.pi * u_xx)**2)

    # backpropagate joint loss, take optimiser step
    loss = lambda1*loss1 + lambda2*loss2 + lambda3*loss3
    loss.backward()
    optimiser.step()
```

PDE

Training cycle:

$$|\mathcal{N}\mathcal{N}(x_i) - u_0(x_i)| \longrightarrow$$

$$|\mathcal{N}\mathcal{N}(-1, t_i) - 0 + \mathcal{N}\mathcal{N}(1, t_i) - 0|^2 \longrightarrow$$

```
for i in range(5000):
    optimiser.zero_grad()

    lambda1, lambda2, lambda3 = 1, 1e-4, 1e-1

    # compute initial data loss
    u = pinn(initial_data_nodes_t)
    loss1 = torch.mean((torch.squeeze(u) - initial_data_values_t)**2)

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss2 = torch.mean((u - torch.zeros(u.shape))**2)

    # compute physics loss
    u = pinn(collocation_data_nodes_t)
    du = torch.autograd.grad(u, collocation_data_nodes_t, \
                              torch.ones([collocation_data_nodes_t.shape[0], 1]), \
                              retain_graph=True, create_graph=True)[0]
    du2 = torch.autograd.grad(du, collocation_data_nodes_t, \
                              torch.ones(collocation_data_nodes_t.shape), \
                              create_graph=True)[0]

    u_x = du[:, [0]]
    u_t = du[:, [1]]
    u_xx = du2[:, [0]]
    loss3 = torch.mean((u_t + u*u_x - .01 / np.pi * u_xx)**2)

    # backpropagate joint loss, take optimiser step
    loss = lambda1*loss1 + lambda2*loss2 + lambda3*loss3
    loss.backward()
    optimiser.step()
```

PDE

Training cycle:

$$|\mathcal{N}\mathcal{N}(x_i) - u_0(x_i)| \longrightarrow$$

$$|\mathcal{N}\mathcal{N}(-1, t_i) - 0 + \mathcal{N}\mathcal{N}(1, t_i) - 0|^2 \longrightarrow$$

$$\partial_t u, \partial_x u, \partial_{xx} u$$

```
for i in range(5000):
    optimiser.zero_grad()

    lambda1, lambda2, lambda3 = 1, 1e-4, 1e-1

    # compute initial data loss
    u = pinn(initial_data_nodes_t)
    loss1 = torch.mean((torch.squeeze(u) - initial_data_values_t)**2)

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss2 = torch.mean((u - torch.zeros(u.shape))**2)

    # compute physics loss
    u = pinn(collocation_data_nodes_t)
    du = torch.autograd.grad(u, collocation_data_nodes_t, \
                             torch.ones([collocation_data_nodes_t.shape[0], 1]), \
                             retain_graph=True, create_graph=True)[0]
    du2 = torch.autograd.grad(du, collocation_data_nodes_t, \
                              torch.ones(collocation_data_nodes_t.shape), \
                              create_graph=True)[0]

    u_x = du[:, [0]]
    u_t = du[:, [1]]
    u_xx = du2[:, [0]]

    loss3 = torch.mean((u_t + u*u_x - .01 / np.pi * u_xx)**2)

    # backpropagate joint loss, take optimiser step
    loss = lambda1*loss1 + lambda2*loss2 + lambda3*loss3
    loss.backward()
    optimiser.step()
```


PDE

Training cycle:

$$|\mathcal{N}\mathcal{N}(x_i) - u_0(x_i)| \longrightarrow$$

$$|\mathcal{N}\mathcal{N}(-1, t_i) - 0 + \mathcal{N}\mathcal{N}(1, t_i) - 0|^2 \longrightarrow$$

$$\partial_t u, \partial_x u, \partial_{xx} u$$

$$|f(u, x_i, t_i)|^2 \longrightarrow$$

```
for i in range(5000):
    optimiser.zero_grad()

    lambda1, lambda2, lambda3 = 1, 1e-4, 1e-1

    # compute initial data loss
    u = pinn(initial_data_nodes_t)
    loss1 = torch.mean((torch.squeeze(u) - initial_data_values_t)**2)

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss2 = torch.mean((u - torch.zeros(u.shape))**2)


    # compute physics loss
    u = pinn(collocation_data_nodes_t)
    du = torch.autograd.grad(u, collocation_data_nodes_t, \
                             torch.ones([collocation_data_nodes_t.shape[0], 1]), \
                             retain_graph=True, create_graph=True)[0]
    du2 = torch.autograd.grad(du, collocation_data_nodes_t, \
                              torch.ones(collocation_data_nodes_t.shape), \
                              create_graph=True)[0]

    u_x = du[:, [0]]
    u_t = du[:, [1]]
    u_xx = du2[:, [0]]
    loss3 = torch.mean((u_t + u*u_x - .01 / np.pi * u_xx)**2)

    # backpropagate joint loss, take optimiser step
    loss = lambda1*loss1 + lambda2*loss2 + lambda3*loss3
    loss.backward()
    optimiser.step()
```



PDE

Training cycle:






PDE

Training cycle:







PDE

Training cycle:



PDE

Training cycle:



ODE

- The displacement of the spring follows the following differential equation:

$$m \frac{d^2}{dt^2} u + \mu \frac{d}{dt} u + ku = 0$$

m is the mass of the oscillator, μ the coefficient of friction and k the spring constant.

Considering the case where the oscillations are damped by friction and initial conditions of the system are $u(t = 0) = 1$, $\frac{d}{dt}u(t = 0) = 0$.

There is an analytical solution for this problem.

https://beltoforion.de/en/harmonic_oscillator/

ODE

- The displacement of the spring follows the following differential equation:

$$m \frac{d^2}{dt^2} u + \mu \frac{d}{dt} u + ku = 0$$

m is the mass of the oscillator, μ the coefficient of friction and k the spring constant.

Considering the case where the oscillations are damped by friction and initial conditions of the system are $u(t = 0) = 1$, $\frac{d}{dt} u(t = 0) = 0$.

There is an analytical solution for this problem.

Use PINNs to solve this ODE

https://beltoforion.de/en/harmonic_oscillator/

ODE

Define the sampling nodes

```
boundary_nodes = np.zeros(1)
collocation_nodes = np.linspace(0,1,30)
test_nodes = np.linspace(0,1,300)

boundary_nodes_t = torch.tensor(boundary_nodes, dtype=torch.float).view(-1,1)
collocation_nodes_t = torch.tensor(collocation_nodes, dtype=torch.float).view(-1,1)
test_nodes_t = torch.tensor(test_nodes, dtype=torch.float).view(-1,1)

boundary_nodes_t.requires_grad = True
collocation_nodes_t.requires_grad = True
test_nodes_t.requires_grad = True
```

- Source: <https://github.com/benmoseley/DLSC-2023>

ODE

Neural network architecture the same as for PDE (except input is dimension 1).

Training cycle:

\mathcal{L}_b



```
for i in range(15001):
    optimiser.zero_grad()

    lambda1, lambda2 = 1e-1, 1e-4

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss1 = (torch.squeeze(u) - 1)**2
    dudt = torch.autograd.grad(u, boundary_nodes_t,\
                                torch.ones_like(u), retain_graph=True, create_graph=True)[0]
    loss2 = (torch.squeeze(dudt) - 0)**2

    # compute physics loss
    u = pinn(collocation_nodes_t)
    dudt = torch.autograd.grad(u, collocation_nodes_t,\
                                torch.ones_like(u), create_graph=True)[0]
    d2udt2 = torch.autograd.grad(dudt, collocation_nodes_t,\
                                   torch.ones_like(dudt), create_graph=True)[0]
    loss3 = torch.mean((d2udt2 + mu*dudt + k*u)**2)

    # backpropagate joint loss, take optimiser step
    loss = loss1 + lambda1*loss2 + lambda2*loss3
    loss.backward()
    optimiser.step()
```

ODE

Neural network architecture the same as for PDE (except input is dimension 1).

Training cycle:

\mathcal{L}_b



```
for i in range(15001):
    optimiser.zero_grad()

    lambda1, lambda2 = 1e-1, 1e-4

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss1 = (torch.squeeze(u) - 1)**2
    dudt = torch.autograd.grad(u, boundary_nodes_t,\
                                torch.ones_like(u), retain_graph=True, create_graph=True)[0]
    loss2 = (torch.squeeze(dudt) - 0)**2

    # compute physics loss
    u = pinn(collocation_nodes_t)
    dudt = torch.autograd.grad(u, collocation_nodes_t,\
                                torch.ones_like(u), create_graph=True)[0]
    d2udt2 = torch.autograd.grad(dudt, collocation_nodes_t,\
                                torch.ones_like(dudt), create_graph=True)[0]
    loss3 = torch.mean((d2udt2 + mu*dudt + k*u)**2)

    # backpropagate joint loss, take optimiser step
    loss = loss1 + lambda1*loss2 + lambda2*loss3
    loss.backward()
    optimiser.step()
```

$\frac{d}{dt}u, \frac{d^2}{dt^2}u$

ODE

Neural network architecture the same as for PDE (except input is dimension 1).

Training cycle:

\mathcal{L}_b

$$\frac{d}{dt}u, \frac{d^2}{dt^2}u$$

\mathcal{L}_f

```
for i in range(15001):
    optimiser.zero_grad()


    lambda1, lambda2 = 1e-1, 1e-4

    # compute boundary loss
    u = pinn(boundary_nodes_t)
    loss1 = (torch.squeeze(u) - 1)**2
    dudt = torch.autograd.grad(u, boundary_nodes_t,\
                                torch.ones_like(u), retain_graph=True, create_graph=True)[0]
    loss2 = (torch.squeeze(dudt) - 0)**2



    # compute physics loss
    u = pinn(collocation_nodes_t)
    dudt = torch.autograd.grad(u, collocation_nodes_t,\
                                torch.ones_like(u), create_graph=True)[0]
    d2udt2 = torch.autograd.grad(dudt, collocation_nodes_t,\
                                torch.ones_like(dudt), create_graph=True)[0]
    loss3 = torch.mean((d2udt2 + mu*dudt + k*u)**2)

    # backpropagate joint loss, take optimiser step
    loss = loss1 + lambda1*loss2 + lambda2*loss3
    loss.backward()
    optimiser.step()
```




ODE



ODE



ODE



ODE

